

Overall Goal:

This notebook takes the pre-trained **Qwen2.5-VL 7B Instruct** model (loaded efficiently using Unsloth) and **fine-tunes** it. The goal of the fine-tuning is likely to teach the model not just to *predict* the answer to visual multiple-choice questions (like in the first notebook), but to also **generate a structured, step-by-step reasoning process** explaining *how* it arrived at the (known) correct answer, outputting this reasoning in a specific JSON format.

The Two-Stage Process:

This notebook executes a two-stage strategy:

1. Stage 1: Generating Structured Training Data (Using the Base Model):

- It starts identically to the first notebook: installing libraries (unsloth, transformers, torch), loading the 4-bit quantized Qwen VLM via Unsloth, configuring it with LoRA adapters (get_peft_model), and loading the appropriate processor.
- It also loads the EXAMS-V dataset and preprocesses the first 200 images (resizing, saving locally to /content/downloaded_images).
- **Crucially**, it then runs an *inference loop* (Cell 10) that is different from the first notebook. Instead of just predicting, it:
 - Loads the *ground truth answers* for the 200 images from the results.json file created during image preprocessing.
 - For each image, it **prompts the model** specifically asking it to generate a detailed step-by-step reasoning process (step_1 to step_4, plus other details like extracted question/choices) that leads to the **provided correct answer**.
 - It parses the model's generated text output (which should ideally be a JSON containing the reasoning).
 - It **enforces** that the correct_answer field in the output JSON matches the ground truth answer loaded earlier.
 - It saves this structured output (image path + model-generated reasoning + verified correct answer) line by line into a new file: **output.jsonl**. This file becomes the training data for the next stage.

2. Stage 2: Fine-Tuning the Model (Using Generated Data):

- It sets up the environment for training by ensuring necessary libraries (trl for SFTTrainer, peft, accelerate, etc.) are installed.
- It **loads the output.jsonl file** created in Stage 1. This file now contains examples where the input is (implicitly) the image and question, and the target output is the detailed JSON reasoning string.

- It converts this loaded data into a Hugging Face Dataset and splits it into training and testing sets.
- It defines a `formatting_func` to structure each data point into the `<s>[INST] <image_path>...</image_path>\nquestion: ... [/INST] {JSON explanation}</s>` format expected by the SFTTrainer (Supervised Fine-tuning Trainer).
- It configures `TrainingArguments` with hyperparameters like learning rate, number of epochs (3), batch size (using gradient accumulation for an effective size of 8), optimization settings, and saving strategies.
- It initializes the SFTTrainer, providing the PEFT-configured model, tokenizer, formatted datasets, and training arguments.
- It **starts the fine-tuning process** (`trainer.train()`). During training, the model learns to generate the detailed JSON reasoning when prompted with an image and question, based on the examples created in Stage 1.
- Finally, it **saves the fine-tuned LoRA adapters** (the small set of weights that were actually trained) and the tokenizer configuration to a new directory (`./vinayak_model_new`).

Key Differences from File 1:

- **Objective:** File 1 was about *predicting* answers with the base model. File 2 is about *generating training data* (reasoning) using the base model and then *fine-tuning* the model on that data.
- **Inference Prompt:** The prompt in File 2 explicitly provides the correct answer and asks for the reasoning *leading to it*.
- **Output File:** File 1 created `output_raw.jsonl` (image path + predicted answer). File 2 creates `output.jsonl` (image path + full JSON reasoning using the *ground truth* answer).
- **Training Step:** File 2 includes a full fine-tuning process using SFTTrainer on the data generated in its first stage.
- **Final Output:** File 2 saves the fine-tuned model adapters, not just predictions.

In essence, this notebook uses the base model as a 'teacher' to generate reasoning explanations for correct answers, and then fine-tunes the model itself to become better at generating such explanations. The "No Output" in the filename simply means the cell outputs (like progress bars, logs) were cleared before saving the notebook file structure. The code itself clearly defines this two-stage generate-then-train workflow.